

Consider the following XML sample intended to model a *directed graph*.

```
<graph>
  <id>g0</id>
  <name>The Graph Name</name>
  <nodes>
    <node>
      <id>a</id>
      <name>A name</name>
    </node>
    ...
    <node>
      <id>e</id>
      <name>E name</name>
    </node>
  </nodes>
  <edges>
    <node>
      <id>e1</id>
      <from>a</from>
      <to>e</id>
      <cost>42</cost>
    </node>
    ..
    <node>
      <id>e5</id>
      <from>a</from>
      <to>a</to>
      <cost>0.42</cost>
    </node>
  </edges>
</graph>
```

Write a program in your favorite language to:

1. Download a text file containing an XML file like the one described.
2. Parse the file to make sure it is syntactically and semantically correct subject to the following restrictions:
  - There must be an `<id>` and `<name>` for the `<graph>`.
  - Assume the `<nodes>` group will *always* come before the `<edges>` group.
  - There must be at least one `<node>` in the `<nodes>` group.
  - All nodes must have different `<id>` tags.

- For every `<edge>`, there must be a single `<from>` tag and a single `<to>` tag, corresponding to nodes that *must* have been defined before.
- The `<cost>` tag is *optional* and can provide an arbitrary non-negative floating point number. If it's not present, then the cost for the `<edge>` defaults to zero.

You'd have chosen an XML parsing library for your favorite programming language to tackle this task. Please explain why you chose the particular library.

3. Propose a normalized SQL schema to model this graphs in PostgreSQL using *standard SQL data types only*.

Please explain briefly each attribute and relationship you propose. You can use SQL documentation facilities and comments in order to do that.

4. Write an SQL query that finds cycles in a given graph, according to the data model you proposed on item (3). You can use standard SQL99 or PL/pgSQL functions.
5. Suppose there's a frontend application where end-users can create queries having this JSON form

```
{
  "queries" : [
    ...
    "paths" : {
      "start" : "a",
      "end"   : "e"
    },
    ...
    "cheapest" : {
      "start" : "a",
      "end"   : "e"
    },
    ...
  ]
}
```

Write a program that receives this structure via standard input, queries the database, and answers via standard output. The possible queries are to be interpreted as follows:

- There can be many interleaved `paths` and `cheapest` queries, but at least *one* query.
- A *single* JSON document should be emitted, holding all the answers.

- Each **paths** query must return a list of possible paths between the **start** and **end** nodes referenced by **id**, if there are paths between them; an empty list if there aren't any paths from **start** to **end**. If the graph has cycles, they are to be ignored.

A sample answer could be:

```
{
  "answers" : [
    ...
    "paths" : {
      "from" : "a",
      "to" : "e",
      "paths" : [ [ "a", "b", "e" ], [ "a", "e" ] ]
    },
    ...
  ]
}
```

assuming there are paths **a->b->e** and **a->e**.

- Each **cheapest** query must return the cheapest path between the **start** and **end** nodes referenced by **id**, if there's at least one path; a false value if there aren't any paths from **start** to **end**. If the graphs has cycles, they are to be ignored.

A sample answer could be:

```
{
  "answers" : [
    ...
    "cheapest" : {
      "from" : "a",
      "to" : "e",
      "path" : [ "a", "e" ]
    },
    ...
    "cheapest" : {
      "from" : "a",
      "to" : "h",
      "path" : false,
    },
    ...
  ]
}
```

assuming  $a \rightarrow e$  is the cheapest path, and there's no path from  $a$  to  $h$ .

You'd have chosen a JSON parsing/serialization library for your favorite programming language to tackle this task. Please explain why you chose the particular libraries.

Please explain how you solved the “are there any paths” and the “cheapest path” problems.